

PrismDB vs Google Cloud Spanner: A Comprehensive Technical Comparison

Whitepaper Version: 1.0 Date: December 2025 Authors: PrismDB Team

Executive Summary

This whitepaper provides an in-depth feature-by-feature comparison between **PrismDB**, a modern embedded analytical database written in Rust, and **Google Cloud Spanner**, Google's globally distributed, strongly consistent relational database. These systems represent fundamentally different approaches to database design—PrismDB focuses on embedded OLAP workloads while Spanner targets globally distributed OLTP with recent HTAP capabilities.

Key Findings:

- PrismDB is an **embedded, single-node OLAP database** while Spanner is a **globally distributed, multi-region OLTP/HTAP database**
- Spanner uses **TrueTime** for global clock synchronization enabling external consistency; PrismDB uses traditional **MVCC** for local transactions
- Both support **columnar storage** and **vectorized execution** for analytical queries
- Spanner offers **99.999% availability** with automatic replication; PrismDB is designed for **local/embedded** use
- Spanner has evolved into a **multi-model database** with graph, vector, and full-text search; PrismDB has similar capabilities on its roadmap
- PrismDB is **open-source** and **free**; Spanner is a **managed cloud service** with usage-based pricing

Table of Contents

1. [Introduction](#)
2. [Architecture Comparison](#)
3. [Deployment Models](#)
4. [Storage Engine](#)
5. [Query Engine](#)
6. [Data Types](#)
7. [SQL Features](#)
8. [Indexing Mechanisms](#)
9. [Transaction Support](#)
10. [Consistency & Replication](#)
11. [Multi-Model Capabilities](#)
12. [Performance & Scalability](#)
13. [Use Cases](#)
14. [Feature Comparison Matrix](#)
15. [Conclusion](#)

1. Introduction

1.1 PrismDB Overview

PrismDB is a high-performance analytical database written in Rust, designed for OLAP workloads. It emphasizes:

- **Embedded deployment:** In-process execution with zero external dependencies
- **Rust-native implementation:** Memory safety without garbage collection
- **Python integration:** First-class bindings via PyO3
- **ACID compliance:** Full transaction support with MVCC
- **Columnar storage:** Optimized for analytical query patterns

Deployment Model: Local/embedded, single-node **License:** MIT (Open Source) **Current Version:** 0.1.0 (Active Development)

1.2 Google Cloud Spanner Overview

Google Cloud Spanner is a fully managed, globally distributed relational database service. Originally developed internally at Google starting in 2007, it combines:

- **Global distribution:** Multi-region deployments with automatic replication
- **Strong consistency:** External consistency via TrueTime
- **Unlimited scale:** Horizontal scaling across nodes and regions
- **High availability:** Up to 99.999% availability SLA
- **Multi-model capabilities:** Relational, graph, vector, and full-text search

Deployment Model: Managed cloud service (Google Cloud) **License:** Proprietary (Cloud Service) **Current Version:** Continuously updated managed service


1.3 Fundamental Differences

Aspect	PrismDB	Google Spanner
Primary Focus	Embedded OLAP	Distributed OLTP/HTAP
Deployment	Local/embedded	Cloud-managed
Scale	Single node	Global distribution
Consistency	Local MVCC	External consistency (TrueTime)
Availability	Application-dependent	99.999% SLA
Cost Model	Free (open source)	Pay-per-use


2. Architecture Comparison

2.1 System Architecture

PrismDB Architecture:



Google Spanner Architecture:



2.2 Component Comparison

Component	PrismDB	Google Spanner
Query Parser	Custom SQL parser	GoogleSQL / PostgreSQL dialects
Query Optimizer	Rule-based + cost-based	Cost-based with statistics
Execution Engine	Vectorized (pull-based)	Vectorized (for columnar engine)
Storage Engine	Block-based columnar	Ressi (PAX) + Columnar Engine
Transaction Manager	Local MVCC	Distributed 2PC + Paxos
Replication	None (single node)	Paxos-based synchronous
Clock Synchronization	System clock	TrueTime (GPS + atomic clocks)

2.3 Design Philosophy

Philosophy	PrismDB	Google Spanner
CAP Theorem	CP (single node)	CP (with TrueTime)
Consistency Priority	Strong (local)	External consistency (global)
Availability Model	Application-managed	Managed HA (99.999%)
Latency Target	Microseconds (local)	Milliseconds (global)
Scale Model	Vertical	Horizontal

3. Deployment Models

3.1 PrismDB Deployment

Embedded/In-Process:

```
import prismdb

# In-memory database
db = prismdb.connect()

# File-based persistent database
db = prismdb.connect('analytics.db')

# Execute queries
result = db.execute("SELECT * FROM sales GROUP BY region")
```

Characteristics:

- Zero infrastructure required
- No network latency
- Single-file database
- Application-lifecycle bound
- No operational overhead

3.2 Google Spanner Deployment

Cloud-Managed Service:

```
from google.cloud import spanner

# Connect to Spanner instance
client = spanner.Client()
instance = client.instance('my-instance')
database = instance.database('my-database')

# Execute queries
with database.snapshot() as snapshot:
    results = snapshot.execute_sql("SELECT * FROM sales")
```

Configuration Options:

Configuration	Description	Availability SLA
Single Region	3 read-write replicas in one region	99.99%
Dual Region	Replicas across 2 regions	99.99%
Multi-Region	Replicas across 3+ regions	99.999%

Editions (as of 2024):

- **Standard:** Core relational capabilities
- **Enterprise:** Multi-model (graph, vector, full-text search), autoscaling

3.3 Deployment Comparison

Aspect	PrismDB	Google Spanner
Setup Time	Seconds	Minutes
Infrastructure	None	Google Cloud

Network Required	No	Yes
Operational Overhead	None	Managed
Geographic Distribution	N/A	Global
Minimum Cost	Free	~\$65/month (1 node)


4. Storage Engine

4.1 Storage Format

PrismDB Storage:

- Block-based columnar storage (256KB blocks)
- Separate files per column within blocks
- Validity masks for NULL tracking
- Column-level statistics (min, max, null count)
- MVCC version chains

Block Structure:



Google Spanner Storage:

- **Ressi Format:** Modern PAX-based columnar format
- Row groups with column-wise layout within blocks
- Supports both point lookups and analytical scans
- **Columnar Engine:** Dedicated columnar storage for analytics
- Automatic compression with high ratios

Spanner Storage Hierarchy:

Universe → Zones → Spanservers → Tablets → Splits → Blocks

4.2 Storage Features

Feature	PrismDB	Google Spanner
Storage Format	Custom columnar	Ressi (PAX) + Columnar
Block Size	256KB	Variable
Row Groups	Configurable	Automatic (splits)
Compression	Dictionary, RLE	Automatic (high ratio)
Column Pruning	Yes	Yes
Late Materialization	Yes	Yes
Storage Tiering	No	Automatic

Encryption at Rest	Planned	Yes (default)
--------------------	---------	---------------

4.3 Data Organization

PrismDB:

- Tables stored as collections of column files
- Optional indexing via B-tree/hash structures
- Single-file database option

Spanner:

- Tables sharded into **splits** by primary key
- Splits distributed across **tablets** on spanservers
- **Interleaved tables** for parent-child co-location
- Automatic split management based on size/load

5. Query Engine

5.1 Execution Model

Aspect	PrismDB	Google Spanner
Execution Style	Vectorized (pull-based)	Row-based + Vectorized (columnar)
Vector Size	2048 tuples	Variable
Parallelism	Morsel-driven	Distributed + local parallel
JIT Compilation	Planned	No
SIMD	Per-function	Yes (columnar engine)

5.2 Query Processing

PrismDB Query Flow:

```
SQL → Tokenizer → Parser → AST → Binder → Logical Plan
      → Optimizer → Physical Plan → Vectorized Executor → Results
```

Spanner Query Flow:

```
SQL → Parser → AST → Analyzer → Distributed Planner
      → Coordinator → Spanserver Execution → Aggregation → Results
```

5.3 Physical Operators

Operator	PrismDB	Google Spanner
Table Scan	Yes	Yes
Index Scan	Yes	Yes
Filter	Yes	Yes
Projection	Yes	Yes
Hash Join	Yes	Yes

Sort-Merge Join	Yes	Yes
Distributed Join	No	Yes
Hash Aggregate	Yes	Yes
Distributed Aggregate	No	Yes
Sort	Yes	Yes
Limit/Top-N	Yes	Yes
Window Functions	Yes	Yes

5.4 Query Optimization

Optimization	PrismDB	Google Spanner
Filter Pushdown	Yes	Yes
Projection Pushdown	Yes	Yes
Constant Folding	Yes	Yes
Join Reordering	Yes	Yes
Distributed Query Planning	No	Yes
Automatic Index Selection	Yes	Yes
Statistics-based Optimization	Basic	Advanced
Query Hints	Basic	Yes (FORCE_INDEX)

6. Data Types

6.1 Scalar Types

Type Category	PrismDB	Google Spanner
Boolean	BOOLEAN	BOOL
Integers	TINYINT, SMALLINT, INTEGER, BIGINT, HUGEINT	INT64 only
FLOATs	FLOAT, DOUBLE	FLOAT32, FLOAT64
Decimals	DECIMAL(p,s)	NUMERIC (precision 38, scale 9)
Strings	VARCHAR, CHAR, TEXT	STRING
Binary	BLOB	BYTES

6.2 Temporal Types

Type	PrismDB	Google Spanner
Date	DATE	DATE
Time	TIME	Not directly (use TIMESTAMP)
Timestamp	TIMESTAMP	TIMESTAMP

Interval	INTERVAL	INTERVAL
Time Zone	Basic	TIMESTAMP handles UTC

6.3 Complex Types

Type	PrismDB	Google Spanner
Arrays	LIST(T), ARRAY	ARRAY
Structs	STRUCT(fields)	STRUCT
Maps	MAP(K, V)	Not native (use STRUCT array)
JSON	JSON	JSON
Enums	ENUM(values)	ENUM
UUID	UUID	Not native (use STRING/BYTES)
Protocol Buffers	No	PROTO
Graph Elements	Planned	GRAPH_ELEMENT, GRAPH_PATH

6.4 Type System Comparison

Spanner Advantages:

- Protocol Buffer (PROTO) native support
- Graph element types for Spanner Graph
- Tight integration with Google ecosystem

PrismDB Advantages:

- Multiple integer sizes (memory efficiency)
- Native UUID type
- Native MAP type
- More flexible decimal precision

7. SQL Features

7.1 SQL Dialects

PrismDB:

- Custom SQL dialect (PostgreSQL-influenced)
- Single dialect

Spanner:

- **GoogleSQL:** Primary dialect, compatible with BigQuery
- **PostgreSQL:** Wire-compatible interface for portability

7.2 DDL Features

Feature	PrismDB	Google Spanner
CREATE TABLE	Yes	Yes
CREATE OR REPLACE	Yes	Limited

ALTER TABLE	Basic	Extensive
DROP TABLE	Yes	Yes
CREATE VIEW	Yes	Yes
MATERIALIZED VIEW	Yes	No (use change streams)
CREATE INDEX	Yes	Yes
INTERLEAVED TABLES	No	Yes (unique feature)
DROP PROTECTION	No	Yes

7.3 DML Features

Feature	PrismDB	Google Spanner
SELECT	Yes	Yes
INSERT	Yes	Yes
UPDATE	Yes	Yes
DELETE	Yes	Yes
MERGE/UPSERT	Planned	INSERT OR UPDATE
Partitioned DML	No	Yes
Batch DML	Via transaction	Yes

7.4 Advanced SQL Features

Feature	PrismDB	Google Spanner
CTEs (WITH clause)	Yes	Yes
Recursive CTEs	In progress	Yes
Window Functions	Full support	Full support
PIVOT/UNPIVOT	Yes	Limited
QUALIFY	Yes	No
ARRAY Functions	Basic	Extensive
JSON Functions	Basic	Extensive
ML Functions	No	Yes (ML.PREDICT)
Graph Queries (GQL)	Planned	Yes (Spanner Graph)

7.5 Join Types

Join Type	PrismDB	Google Spanner
INNER JOIN	Yes	Yes
LEFT/RIGHT JOIN	Yes	Yes

FULL OUTER JOIN	Yes	Yes
CROSS JOIN	Yes	Yes
SEMI JOIN	Yes	Yes
ANTI JOIN	Yes	Yes
Hash Join	Yes	Yes
Distributed Join	No	Yes

8. Indexing Mechanisms

8.1 Index Types

Index Type	PrismDB	Google Spanner
Primary Key Index	Yes	Yes (automatic)
B-Tree Index	Yes	No (uses primary key sorting)
Hash Index	Yes	No
Secondary Index	Yes	Yes
Interleaved Index	No	Yes
Search Index	Planned	Yes (full-text)
Vector Index	Planned	Yes (ANN/KNN)

8.2 Spanner Unique Index Features

Interleaved Indexes: Spanner can interleave indexes with parent tables for co-location:

```
CREATE INDEX AlbumsByTitle ON Albums(Title)
INTERLEAVE IN Singers;
```

STORING Clause: Store additional columns in the index to avoid base table lookups:

```
CREATE INDEX AlbumsByTitle ON Albums(Title)
STORING (ReleaseDate, Genre);
```

8.3 Index Usage

Feature	PrismDB	Google Spanner
Automatic Index Selection	Yes	Yes
Index Hints	Basic	Yes (FORCE_INDEX)
Covering Indexes	No	Yes (STORING)
Partial Indexes	Planned	No
Expression Indexes	No	No

Index-Only Scans	Limited	Yes
------------------	---------	-----

9. Transaction Support

9.1 ACID Properties

Property	PrismDB	Google Spanner
Atomicity	Full	Full (distributed)
Consistency	Full	Full (external consistency)
Isolation	Multiple levels	Serializable / Repeatable Read
Durability	WAL-based	Synchronous replication

9.2 Isolation Levels

PrismDB Isolation Levels:

- Read Uncommitted
- Read Committed
- Repeatable Read (default)
- Serializable

Spanner Isolation Levels:

- Serializable (default, with external consistency)
- Repeatable Read (preview, optimistic locking)

9.3 Transaction Types

Transaction Type	PrismDB	Google Spanner
Read-Write	Yes	Yes
Read-Only	Yes	Yes (lock-free)
Single Statement	Yes	Yes
Multi-Statement	Yes	Yes
Distributed	No	Yes
Partitioned DML	No	Yes

9.4 Concurrency Control

Aspect	PrismDB	Google Spanner
Mechanism	MVCC	MVCC + 2PC + TrueTime
Lock Granularity	Row-level	Cell-level
Deadlock Detection	Yes	Yes (distributed)
Concurrent Writers	Yes (row-level)	Yes (cell-level)
Read-Write Conflict	MVCC isolation	Pessimistic locking

9.5 TrueTime: Spanner's Innovation

Spanner's most significant innovation is **TrueTime**, a globally synchronized clock service:

TrueTime API:

- TT.now() → returns interval [earliest, latest]
- TT.after(t) → true if t has definitely passed
- TT.before(t) → true if t has definitely not arrived

Uncertainty bound: ~7ms (GPS + atomic clocks)

External Consistency: If transaction T1 commits before T2 starts (in real time), then T1's commit timestamp < T2's commit timestamp. This guarantee is impossible without synchronized clocks.

10. Consistency & Replication

10.1 Consistency Models


Model	PrismDB	Google Spanner
Local Consistency	Strong	Strong
Distributed Consistency	N/A	External consistency
Read Consistency	Snapshot	Snapshot + bounded staleness
Write Consistency	Immediate	Synchronous replication

10.2 Replication

Aspect	PrismDB	Google Spanner
Replication Type	None	Synchronous (Paxos)
Replicas per Node	1	3+ (automatic)
Cross-Region	No	Yes
Consensus Protocol	N/A	Paxos
Leader Election	N/A	Automatic
Failover	Manual	Automatic

10.3 Spanner Replication Architecture

Paxos Group (per split):



10.4 Availability Comparison

Metric	PrismDB	Google Spanner
Single Node Availability	Application-dependent	N/A (always distributed)
Regional Availability	N/A	99.99%
Multi-Region Availability	N/A	99.999%
RTO (Recovery Time)	Application-dependent	Seconds
RPO (Data Loss)	WAL-dependent	Zero (synchronous)

11. Multi-Model Capabilities

11.1 Current Multi-Model Support

Capability	PrismDB	Google Spanner
Relational	Yes	Yes
Key-Value	Via tables	Yes
Document (JSON)	Basic	Yes
Graph	Planned	Yes (Spanner Graph)
Vector Search	Planned	Yes (ANN/KNN)
Full-Text Search	No	Yes
Time Series	Via SQL	Via SQL

11.2 Spanner Graph (2024)

Spanner Graph adds native graph query capabilities:

```
-- Create a property graph
CREATE PROPERTY GRAPH FinGraph
  NODE TABLES (
    Account LABEL Account,
    Person LABEL Person
  )
  EDGE TABLES (
    Transfers SOURCE KEY (from_id) REFERENCES Account
      DESTINATION KEY (to_id) REFERENCES Account
      LABEL Transfers
  );
-- Query with GQL
GRAPH FinGraph
MATCH (p:Person)-[:Owns]->(a:Account)-[t:Transfers]->(b:Account)
WHERE t.amount > 1000
RETURN p.name, SUM(t.amount);
```

11.3 Spanner Vector Search

```
-- Create vector index
CREATE VECTOR INDEX EmbeddingIndex
ON Products(embedding)
OPTIONS (distance_type = 'COSINE', tree_depth = 2);

-- Similarity search
SELECT product_name, embedding
FROM Products
ORDER BY COSINE_DISTANCE(embedding, @query_embedding)
LIMIT 10;
```

11.4 PrismDB HTAP Roadmap

PrismDB's planned multi-model capabilities (from HTAP Technical Design):

- **Document Store:** BSON encoding, hybrid hot/cold storage
- **Vector Database:** HNSW index, SIMD-accelerated distance calculations
- **Graph Database:** Property graph model, CSR representation, Cypher-like queries

12. Performance & Scalability

12.1 Scale Characteristics

Dimension	PrismDB	Google Spanner
Max Data Size	Node memory/disk	Virtually unlimited
Max Throughput	Single node	2+ billion requests/sec (Google)
Horizontal Scaling	No	Automatic
Vertical Scaling	Yes	Yes (node size)
Data Under Management	GB-TB	6+ exabytes (Google)

12.2 Latency Characteristics

Operation	PrismDB	Google Spanner
Point Read	Microseconds	Single-digit milliseconds
Range Scan	Milliseconds	Milliseconds
Write (single row)	Microseconds	5-10 milliseconds
Distributed Write	N/A	10-50 milliseconds
Analytical Query	Seconds	Seconds (up to 200x faster with columnar)

12.3 Spanner Columnar Engine Performance

The columnar engine (2024) provides significant analytical acceleration:

Improvement	Metric
Scan Performance	Up to 200x faster
Analytical Queries	10-40x improvement

Compression	Higher ratios (columnar layout)
Vectorized Execution	Cache-friendly, SIMD-enabled

12.4 Resource Requirements

Resource	PrismDB	Google Spanner
Minimum Memory	~10 MB	N/A (managed)
Minimum Storage	0 (in-memory)	N/A (managed)
Minimum Nodes	1 (embedded)	1 (but 3 replicas)
Network	None required	Required

13. Use Cases

13.1 Where PrismDB Excels

1. Embedded Analytics

- In-application data processing
- Python data science workflows
- Jupyter notebook integration
- Local file analysis

2. Edge/Offline Scenarios

- IoT edge devices
- Mobile applications
- Disconnected environments
- Low-latency requirements

3. Development & Prototyping

- Zero-infrastructure setup
- Fast iteration cycles
- Learning and education

4. Cost-Sensitive Deployments

- No cloud costs
- Predictable resource usage
- Small to medium datasets

13.2 Where Google Spanner Excels

1. Global Applications

- Multi-region deployments
- Global user base
- Low-latency worldwide access

2. Mission-Critical Workloads

- 99.999% availability requirements
- Zero data loss tolerance
- Financial transactions

3. Massive Scale

- Petabyte+ datasets
- Billions of transactions/day
- Horizontal scaling needs

4. Multi-Model Requirements

- Graph queries on relational data
- Vector search for AI/ML
- Full-text search integration

5. Enterprise Compliance

- Managed security
- Encryption by default
- Audit logging

13.3 Decision Matrix

Requirement	Recommendation
Global distribution	Spanner
Zero infrastructure	PrismDB
99.999% availability	Spanner
Free/open source	PrismDB
Embedded in application	PrismDB
Multi-region consistency	Spanner
Local OLAP processing	PrismDB
Graph + Relational queries	Spanner
Offline capability	PrismDB
Managed operations	Spanner
Python data science	PrismDB
Petabyte scale	Spanner

14. Feature Comparison Matrix

14.1 Core Database Features

Feature	PrismDB	Google Spanner
SQL Support	✓	✓
ACID Transactions	✓	✓
Columnar Storage	✓	✓
Vectorized Execution	✓	✓ (columnar engine)
Distributed Transactions	✗	✓
Global Distribution	✗	✓

Automatic Replication	✗	✓
High Availability	✗	✓ (99.999%)
TrueTime/External Consistency	✗	✓

14.2 SQL Features

Feature	PrismDB	Google Spanner
CTEs	✓	✓
Window Functions	✓	✓
Recursive CTEs	⚠ Partial	✓
PIVOT/UNPIVOT	✓	⚠ Limited
QUALIFY	✓	✗
JSON Functions	⚠ Basic	✓
ML Functions	✗	✓
Graph Queries (GQL)	➡ SOON Planned	✓

14.3 Storage & Indexing

Feature	PrismDB	Google Spanner
B-Tree Index	✓	✗
Secondary Index	✓	✓
Interleaved Tables/Indexes	✗	✓
Covering Indexes (STORING)	✗	✓
Full-Text Search Index	➡ SOON Planned	✓
Vector Index	➡ SOON Planned	✓
Automatic Compression	⚠ Basic	✓

14.4 Multi-Model Capabilities

Feature	PrismDB	Google Spanner
Relational	✓	✓
Key-Value	⚠ Via tables	✓
Document/JSON	⚠ Basic	✓
Graph Database	➡ SOON Planned	✓
Vector Search	➡ SOON Planned	✓
Full-Text Search	✗	✓

14.5 Operations & Management

Feature	PrismDB	Google Spanner
Zero Infrastructure	✓	✗
Managed Service	✗	✓
Automatic Backups	✗	✓
Point-in-Time Recovery	➡️ <small>SOON</small> Planned	✓
Schema Change Protection	✗	✓
Encryption at Rest	➡️ <small>SOON</small> Planned	✓
IAM Integration	✗	✓

14.6 Legend

- ✓ Fully supported
 - ⚠ Partially supported
 - ➡️ SOON Planned/In development
 - ✗ Not supported
-

15. Conclusion

15.1 Summary

PrismDB and Google Cloud Spanner represent opposite ends of the database spectrum:

Dimension	PrismDB	Google Spanner
Philosophy	Simple, embedded	Global, distributed
Scale	Single node	Planet-scale
Deployment	Zero infrastructure	Managed cloud
Cost	Free	\$65+/month
Consistency	Local MVCC	External consistency
Availability	Application-dependent	99.999% SLA
Primary Workload	OLAP	OLTP + HTAP

15.2 When to Choose PrismDB

- **Zero infrastructure** is required
- **Embedded analytics** in applications
- **Cost-sensitive** deployments
- **Offline/edge** scenarios
- **Python/data science** workflows
- **Learning and prototyping**
- **Single-node** workloads are sufficient

15.3 When to Choose Google Spanner

- **Global distribution** is required
- **Mission-critical** availability (99.999%)
- **Petabyte-scale** data

- **Strong global consistency** is mandatory
- **Multi-model** (graph + vector + relational) needed
- **Managed operations** preferred
- **Enterprise compliance** requirements

15.4 Complementary Usage

The two databases can be complementary:

1. **Development:** Use PrismDB for local development and testing
2. **Production:** Deploy to Spanner for global, mission-critical workloads
3. **Data Science:** Use PrismDB for local analytics, Spanner for production data
4. **Edge + Cloud:** PrismDB at the edge, Spanner in the cloud

15.5 Future Outlook

PrismDB Roadmap:

- HTAP capabilities (document store, vector DB, graph DB)
- Enhanced compression (LZ4, ZSTD)
- Distributed query execution (long-term)

Spanner Trajectory:

- Continued multi-model expansion
- Deeper AI/ML integration
- Performance improvements (columnar engine)
- New editions and pricing models

References

1. Google Cloud Spanner Documentation - <https://cloud.google.com/spanner/docs>
2. Spanner: TrueTime and External Consistency - <https://cloud.google.com/spanner/docs/true-time-external-consistency>
3. Spanner Innovations in 2024 - <https://cloud.google.com/blog/products/databases/spanner-innovations-in-2024>
4. Spanner Columnar Engine - <https://cloud.google.com/spanner/docs/columnar-engine>
5. Spanner Graph - <https://cloud.google.com/products/spanner/graph>
6. Spanner Data Types - <https://cloud.google.com/spanner/docs/reference/standard-sql/data-types>
7. Spanner Transactions - <https://cloud.google.com/spanner/docs/transactions>
8. Spanner Secondary Indexes - <https://cloud.google.com/spanner/docs/secondary-indexes>
9. PrismDB Architecture Documentation
10. "Spanner: Google's Globally-Distributed Database" (Corbett et al., OSDI 2012)

Document Version: 1.0 **Last Updated:** December 2025 **License:** MIT

This whitepaper was generated based on analysis of PrismDB source code (version 0.1.0) and publicly available Google Cloud Spanner documentation as of December 2025.